

Lecture 03: Low-Level Software

August 26, 2024

Overview

- ▶ High-level software
- ▶ Low-level aspects
- ▶ High-to-low mapping
- ▶ Assembly languages
- ▶ Compilers and execution

High-Level Perspectives

Humans and computers “think” about software differently. High-level languages allow humans to reason about software more clearly.

- ▶ Program structure
 - ▶ Files, namespaces, modules, classes, etc.
- ▶ Data management
 - ▶ Data structures, memory allocation, scopes, etc.
- ▶ Basic control flow
 - ▶ blocks, loops, conditionals, etc.

It's best to understand the software engineering process and tools.

High-Level Perspectives

Program Structure

Relative advantages and disadvantages

High-Level Perspectives

Program Structure

Relative advantages and disadvantages

- ⊕ Make large complex software manageable
- ⊕ Provides alternative to “spaghetti code”
- ⊕ Breaks program into “units”
 - ▶ Mental tracking, division of effort
- ⊖ Tends to be lost in compilation and linking
 - ▶ units combined
 - ▶ functions inlined
 - ▶ redundancies removed
- ⊖ Must be reconstructed, at least in part, by a reverser

High-Level Perspectives

Program Structure

“Encapsulation” an abstraction of a concept and its related functions into a usable unit of software, which hides the implementation details.

- ▶ Happens at various levels – “nesting of boxes”
- ▶ Each thoroughly tested
- ▶ Each having a well-defined interface

A reverser must have some notion of the program structure in order to navigate to its interesting parts.

High-Level Perspectives

Modules

Largest program unit other than the program itself

Static libraries: `.obj` `.o` `.a`

- ▶ a collection of source files built together
- ▶ typically represent a feature or area of functionality
- ▶ often third-party developed
- ➖ difficult to isolate while reversing

High-Level Perspectives

Modules

Dynamic libraries: `.dll` `.so`

- ▶ like static, but not embedded

High-Level Perspectives

Modules

Dynamic libraries: .dll .so

- ▶ like static, but not embedded
- + allows upgrading components independently
- requires consistent interface
- “DLL hell”
- + easy to isolate while reversing (separate file)
- + exposes interface between components

High-Level Perspectives

Common Code Constructs

Procedures

- ▶ most fundamental unit of software
- ▶ a unit of code with a well-defined purpose
- ▶ can be invoked by other areas
- ▶ can take input and output: “function”
- ▶ most common form of encapsulation

High-Level Perspectives

Common Code Constructs

Objects

- ▶ object-oriented design – supplant procedure as unit of reason
- ▶ probably most popular design methodology
- ▶ contain both data and code
- ▶ code typically manipulates data within the object
- ▶ data typically hidden from external components
- ▶ “clients” interact with objects via an interface

inheritance allows a specific implementation to borrow generic components and interfaces

- ▶ facilitates polymorphism

High-Level Perspectives

Common Code Constructs

Aspects¹

- ▶ permits organization of code into concerns
- ▶ some concerns “cross-cut” the class structure (e.g., logging)
- ▶ usually involves joining “advice” to other code
“When you would ordinarily execute `f()`, also execute this”

Implementations

- ▶ Weave code at the source or bytecode level
- ▶ Or change interpreter to implement AOP

AspectJ extension to Java

Partial classes in C# (not quite)

¹Wikipedia: Aspect Oriented Programming

High-Level Perspectives

Data Management

All programs (and their component operations):

- ▶ take input,
- ▶ use temporary storage, and
- ▶ produce output.

Like code, the high-level constructs are translated to lower-level, causing some information loss.

High-Level Perspectives

Data Management

Data Structures

High-Level Perspectives

Data Management

Data Structures

Variable a named storage location of a type

- ▶ declared with some scope, defining its access and storage
- ▶ name usually replaced by address

High-Level Perspectives

Data Management

Data Structures

Variable a named storage location of a type

- ▶ declared with some scope, defining its access and storage
- ▶ name usually replaced by address

User-defined Structure a group of related fields each of its own type

- ▶ usually stored and handled together
 - ▶ reversing structures is critical to understanding a program
- record incremental details as you go

High-Level Perspectives

Data Management

Data Structures

Collection a group of fields all of the same type

- ▶ come in different flavors and implementations:
 - ▶ expected size
 - ▶ preserves order / sorted
 - ▶ requires uniqueness
 - ▶ frequency/location of searches, insertions, deletions

High-Level Perspectives

Collections

Arrays

- ▶ most basic - items placed sequentially
- ▶ referenced by index
- ▶ can be multidimensional

High-Level Perspectives

Collections

Arrays

- ▶ most basic - items placed sequentially
- ▶ referenced by index
- ▶ can be multidimensional
- + fastest access – direct support on x86
- + most intuitive
- difficult to add and remove at the middle
especially inefficient for large lists

High-Level Perspectives

Collections

Linked Lists

- ▶ items placed individually anywhere
- ▶ each item contains a reference or “link” to the next
- ▶ may also have a link to the previous

High-Level Perspectives

Collections

Linked Lists

- ▶ items placed individually anywhere
- ▶ each item contains a reference or “link” to the next
- ▶ may also have a link to the previous
- + fast insertion and removal
- access by index is slow
- uses more memory for links

High-Level Perspectives

Collections

Trees

- ▶ many kinds (binary, AVL, RB, etc.)
- ▶ items placed individually anywhere
- ▶ items linked in hierarchy rather than in sequence

High-Level Perspectives

Collections

Trees

- ▶ many kinds (binary, AVL, RB, etc.)
- ▶ items placed individually anywhere
- ▶ items linked in hierarchy rather than in sequence
- ⊃ like linked list
- − uses more memory (at least two links)
- + fast searching (implicitly sorted)
- − does not preserve order

High-Level Perspectives

Control Flow

High-level Conditional blocks and loops

Low-level Jumps / conditional jumps

- ▶ Some processors support conditional operations (ARM)
- ▶ Some processors support repeated operations (x86)

Conditional blocks

- ▶ in C: `if`, `else if`, and `else`
- ▶ Conditional controls whether / which block executes

High-Level Perspectives

Control Flow

Switch blocks

- ▶ in C: switch, case, default, and break
- ▶ Takes an input and chooses block(s) to execute
- ▶ Sometimes a lookup table, sometimes converted to if/else

Loops

- ▶ in C: for, do, while, and break
- ▶ Repeated block
- ▶ Interrupted when a condition is met or unmet
- ▶ Usually includes a counter

High-Level Perspectives

Languages

- ▶ Abstract the machine details
- ▶ Assembly simply isn't suitable for complex software
- ▶ Balance simplicity and flexibility
 - ▶ do much with little direction
 - ▶ preclude few operations
- ▶ Reverser is exposed to more details than developer
- ▶ Must seek a high-level understanding
 - ▶ Know the language(s) used
 - ▶ Be familiar with its abstractions

High-Level Perspectives

Languages

C

- ▶ Probably lowest of the high-level languages
- ▶ Arrays have no bounds checking
- ▶ Has data and control structures
- ▶ Is compiled
- ▶ Can be cross-platform at the source level
- ▶ Made for UNIX, but also used by Windows
- ▶ Can be very efficient – little overhead
- ▶ Also very easy to reverse

High-Level Perspectives

Languages

C++

- ▶ Extension to C for OOP
- ▶ Mostly a superset of C
- ▶ Introduces classes – structs with code components and inheritance
- ▶ Supports polymorphism via virtual methods
- ▶ Similar to reversing C – emphasis on determining:
 - ▶ class hierarchy
 - ▶ virtual method calls
 - ▶ constructor calls, etc.

High-Level Perspectives

Languages

Java

- ▶ Compiled to bytecode rather than native machine code
 - ▶ Bytecodes like assembly, but interpreted
- ▶ Cross-platform at the bytecode level – platform must have a JVM
- ▶ Reversing process is very different
- ▶ Easy to decompile with very high accuracy
 - ▶ Binaries tend to be obfuscated as a countermeasure

High-Level Perspectives

Languages

C#

- ▶ Competitor to Java from Microsoft
- ▶ Developed to address shortcomings of C++
- ▶ Bytecodes are called “Microsoft Intermediate Language” or MSIL
- ▶ Cross-platform at the MSIL level – platform must have a .NET VM (Common Language Runtime)
 - ▶ The Microsoft runtime API was limited to Windows
- ▶ Reversing is similar to Java
 - ▶ Binaries often obfuscated

Low-Level Perspectives

Data Management

```
int Multiply(int x, int y) {  
    int z;  
    z = x * y;  
    return z;  
}
```

- ▶ save machine state
- ▶ allocate memory
- ▶ load parameters
- ▶ perform operations (storing intermediate and final values)
- ▶ restore machine state
- ▶ communicate z to the caller
- ▶ return control

Low-Level Perspectives

Data Management

Registers

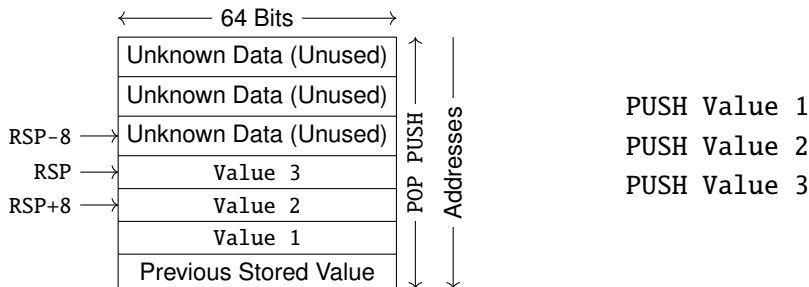
- ▶ A CPU's closest working space
- ▶ To avoid constant RAM access
- ▶ Little or no penalty for access
- ▶ Usually fairly few
- ▶ Managing registers and RAM is not automatic
- ▶ Reverse must determine purpose of register in context

Low-Level Perspectives

Data Management

The Stack

- ▶ Secondary storage area for short-term information
- ▶ Special area in RAM
- ▶ Typically, stacks grow backward (into lower addresses)



Low-Level Perspectives

Data Management

Stack Usage

Low-Level Perspectives

Data Management

Stack Usage

Temporarily-saved register values if a function needs a (callee-save) register, it must save and restore its original value

Local variables if a variable doesn't fit into registers or has operations requiring RAM storage

Function parameters values passed as parameters are usually placed in registers or on the stack

Return addresses the address following a call must be saved to resume later

Low-Level Perspectives

Data Management

Heaps

- ▶ Managed region of dynamically allocatable memory
- ▶ Program requests block of given size, heap returns pointer
- ▶ Typically implemented by library or OS
- ▶ Could be custom implementation
 - ▶ Reversers should locate heaps and management routines
 - ▶ Returned block can be traced
 - ▶ Size of block is known from parameter

Low-Level Perspectives

Data Management

Executable Data Sections

- ▶ Usually named `.data` or `.rodata`
- ▶ Used to store globals and pre-initialized data
- ▶ e.g., C string literals
- ▶ e.g., hard-coded addresses emitted by compiler

Low-Level Perspectives

Control Flow

Control Flow

- ▶ Much more friendly in high-level
- ▶ Conditional statements implemented using primitives
- ▶ Implementations depend on hardware
- ▶ Our discussions will use x86

Compilation

- ▶ Most software is compiled
- ▶ Compilers tend to generate non-intuitive code
- ▶ Primary job is code translation
 - ▶ Usually text source to machine binary
 - ▶ Often apply optimizations
 - ▶ Output not meant for human consumption

Compilation

Architecture

Front end

input: source code

output: intermediate representation

Intermediate Representations

- ▶ used to communicate program from stage to stage
- ▶ sensitive to the purpose of the stage
- ▶ may be closer to high or low level

Compilation

Architecture

Front end

input: source code

lexical analysis groups characters into tokens

syntax analysis checks for correct grammatical structure

semantic analysis derives meaning from structure

output: intermediate representation

Intermediate Representations

- ▶ used to communicate program from stage to stage
- ▶ sensitive to the purpose of the stage
- ▶ may be closer to high or low level

Compilation

Architecture

Optimizer

- ▶ A huge reason to understand compilation
- ▶ Goal is to create most performant, smallest, etc. code

Compilation

Architecture

Optimizer

- ▶ A huge reason to understand compilation
- ▶ Goal is to create most performant, smallest, etc. code

Question: What other goals?

Compilation

Architecture

Optimizer

- ▶ A huge reason to understand compilation
- ▶ Goal is to create most performant, smallest, etc. code

Question: What other goals?

Generic operate on IR

Target-specific operate on machine code

Compilation

Optimizations

Code Structure

loop unrolling produces larger code but reduces branching

switches implemented as table or tree

post-test loops save a conditional branch

Redundancy Elimination

- ▶ eliminating human errors
- ▶ eliminating repeated computations (or pulling outside of a loop)
- ▶ eliminating overwritten and/or unused values
- ▶ caching pointers (e.g., in array indexing)

Compilation

Architecture

Back end (“code generator”)

- ▶ converts IR to target-specific code
- ▶ May perform platform-specific optimizations
(where most annoying transformations take place)

Instruction selection translate IR to machine

Register allocation decide which vars go in hardware registers

Instruction scheduling order to maximize parallelism

Compilation

Architecture

Back end (“code generator”)

- ▶ converts IR to target-specific code
- ▶ May perform platform-specific optimizations
(where most annoying transformations take place)

Instruction selection translate IR to machine

Register allocation decide which vars go in hardware registers

Instruction scheduling order to maximize parallelism

Question: Why separate front and back ends?

Execution

Execution environment The machine that executes a program, its API, libraries, etc.

Execution

Software Execution Environments (VMs)

- ▶ execute an IR program, e.g., Java bytecodes
- ▶ each VM is implemented for a specific machine
- ▶ the IR is common for all implementations of the VM
- ▶ e.g., Java (JVM), and .NET (CLR)

Execution

Software Execution Environments (VMs)

- ▶ execute an IR program, e.g., Java bytecodes
 - ▶ each VM is implemented for a specific machine
 - ▶ the IR is common for all implementations of the VM
 - ▶ e.g., Java (JVM), and .NET (CLR)
- | | | |
|---|----------------------------------|-------------------------|
| + | platform isolation | |
| + | enhanced functionality | ▶ garbage collection |
| - | decreased performance | ▶ runtime type safety |
| - | limited access to low-level APIs | ▶ checked memory access |
| - | limited access to OS features | |

Execution

Software Execution Environments (VMs)

Interpreters

- ▶ original approach to VMs
- ▶ everything is decoded and executed by a program
- ▶ “registers” mapped to memory in interpreter
- performance is very slow
 - ▶ every instruction separately decoded and executed
 - ▶ requires dozens of machine instructions per bytecode instruction

Execution

Software Execution Environments (VMs)

Just-in-Time (JIT) Compilers

- ▶ a compiler back end at run time
- ▶ blocks of bytecode translated to machine code
- ▶ instructions fed to native CPU

Execution

Software Execution Environments (VMs)

Just-in-Time (JIT) Compilers

- ▶ a compiler back end at run time
- ▶ blocks of bytecode translated to machine code
- ▶ instructions fed to native CPU

Question: Why not translate whole program at load?

Execution

Software Execution Environments (VMs)

Just-in-Time (JIT) Compilers

- ▶ a compiler back end at run time
- ▶ blocks of bytecode translated to machine code
- ▶ instructions fed to native CPU

Question: Why not translate whole program at load?

- ▶ Run-time link/loading of modules
- ▶ Translation of polymorphic calls

Execution

Software Execution Environments (VMs)

Reversing Strategies

Execution

Software Execution Environments (VMs)

Reversing Strategies

- ▶ take advantage of metadata (class names, method parameters, variable types, etc.)
- ▶ use a decompiler
 - ⊕ very easy and effective
 - ⊖ motivate equally effective obfuscation techniques
- ▶ bytecodes still reveal higher-level info

Execution

Hardware Execution Environments

processor digital circuits controlled by machine code

simple loop decode instruction, activate circuits

parallelism physical limitations, desire for speed

backward compatibility old programs on new hardware

data dependencies limit parallelism

multiple scheduling execute when input data finishes

Execution

Intel Haswell²

- ▶ latest (circa 2013) Intel microarchitecture used in Core i7
- ▶ Micro-operation cache of 1.5K operations
- ▶ 14- to 19-stage instruction pipeline
- ▶ 4 ALUs per core, 3 AGUs, 2 branch predictors
- ▶ HyperThreading
- ▶ integrated GPU
- ▶ typically 4+ cores
- ▶ dynamically partitioned decode cache in HyperThreading
- ▶ support for NUMA (non-uniform memory access)

²http://en.wikipedia.org/wiki/Haswell_%28microarchitecture%29

Execution

μ -ops

- ▶ x86 instructions translated to μ -instructions executed internally
- ▶ core fairly primitive, most x86 logic in decoder
- ▶ core significantly improved based on now-defunct Itanium
- ▶ microcodes loaded from ROM (cached for performance)
- ▶ microcodes can be patched

Execution

Pipelines³

Instructions broken into stages within a pipeline

Front end decodes each instruction in μ -ops

Out-of-order core executes instructions when resources are available

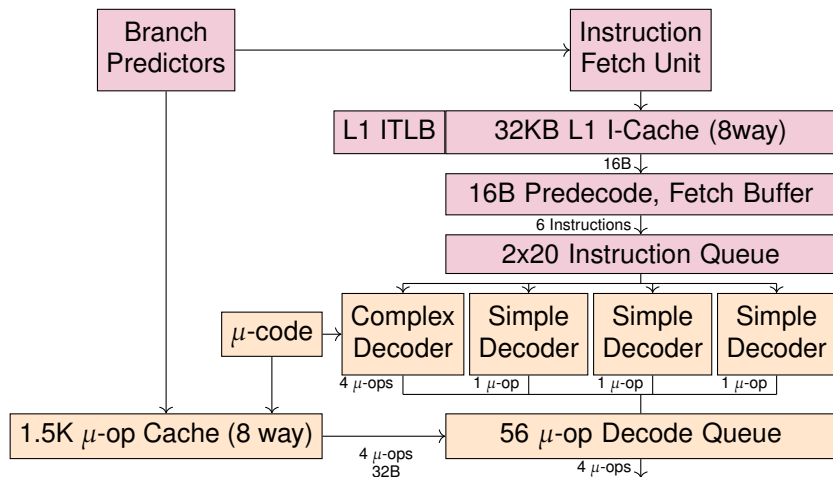
- ▶ can dispatch up to 8 μ -ops per cycle
- ▶ has 8 execution ports
- ▶ can use pipeline forwarding
- ▶ can re-order instructions preserving data dependencies

Retirement section ensure proper ordering

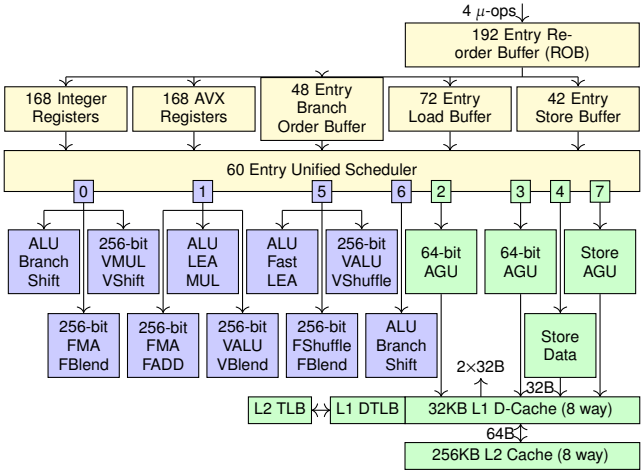
- ▶ 4 μ -ops per cycle per core

³<http://www.realworldtech.com/haswell-cpu/2/>

Execution



Execution



Execution

Branch Prediction

Knowing which instructions to load after conditional branches

Two approaches

- ▶ Stall pipeline until result of branch can be known
 - ▶ Attempt to predict and take appropriate path
 - ▶ an incorrect prediction can be very wasteful
- heuristic e.g., forward vs. backward (loops)
- tracing record previous results

Features affect how optimizers may behave